

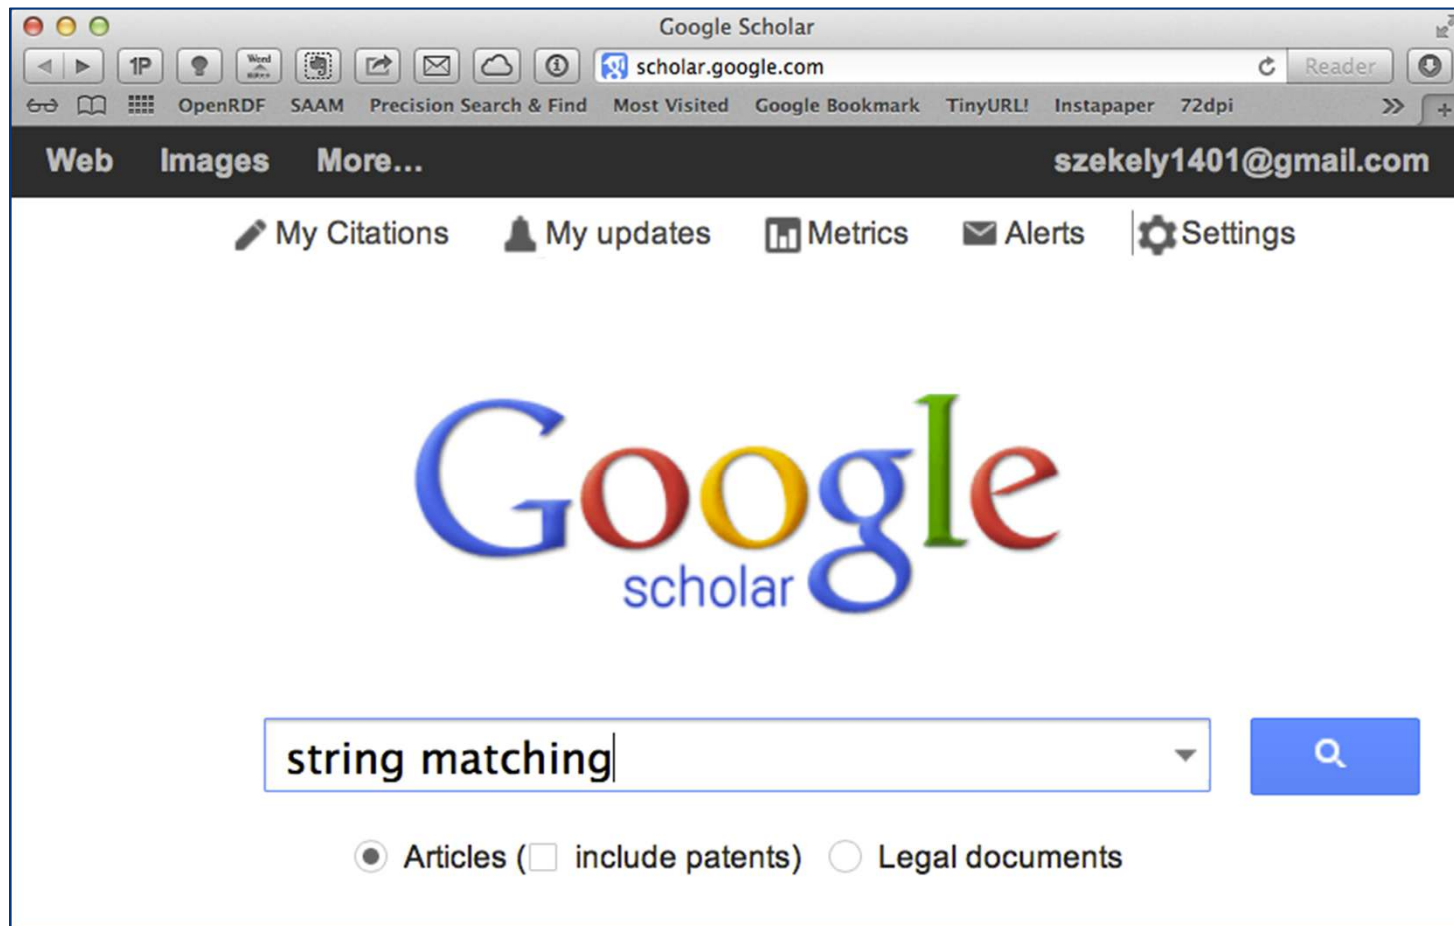
# String Matching

**Pedro Szekely & Craig Knoblock**  
**University of Southern California**

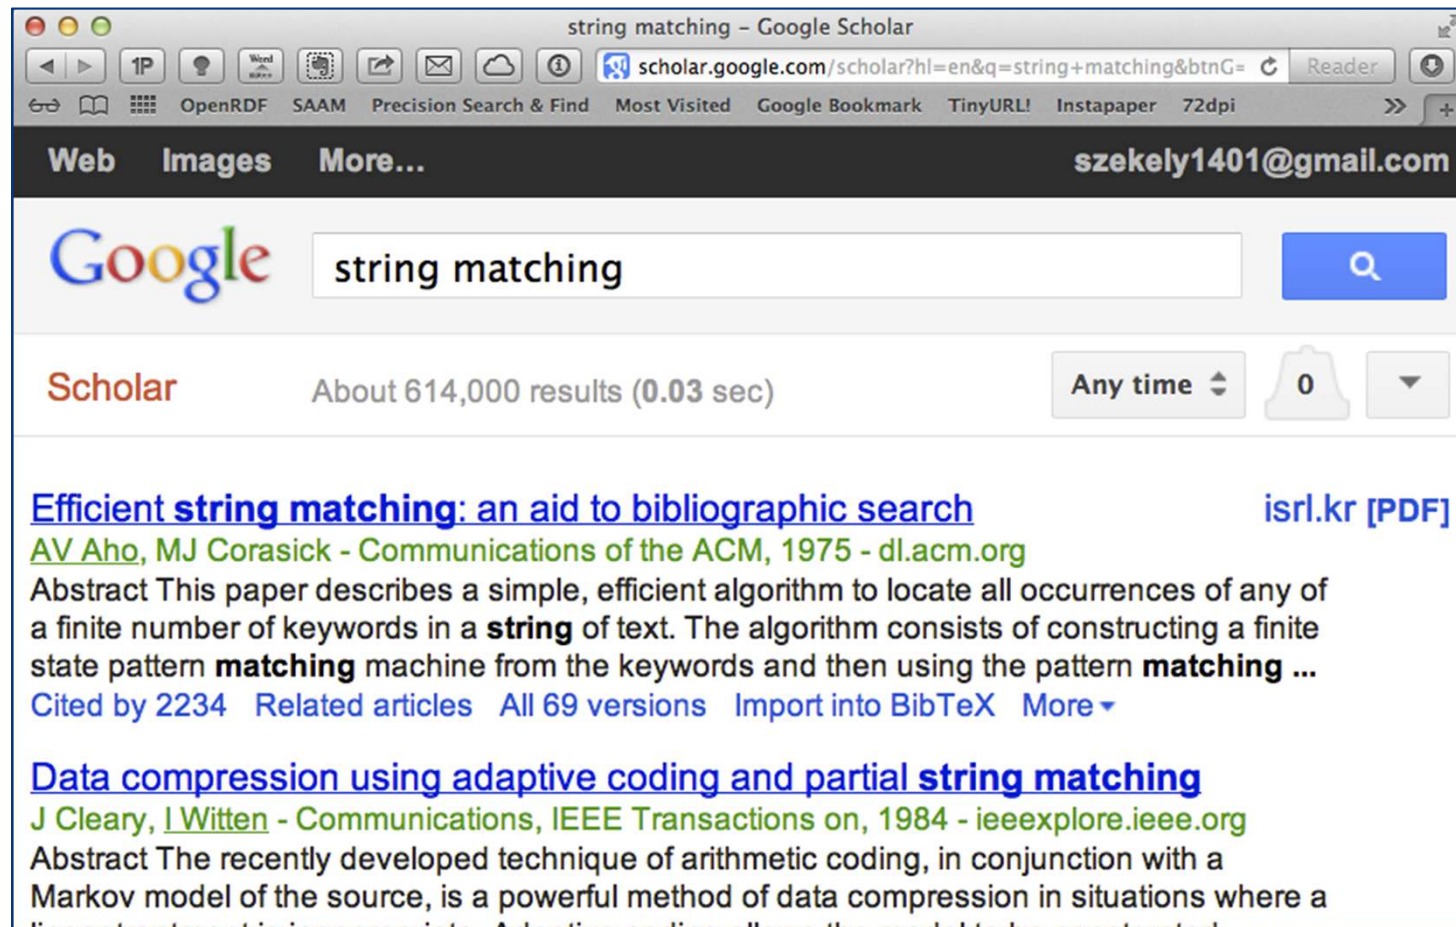
# Isn't the Problem Solved?

`String.equalsIgnoreCase(String x)`

# How Many Publications?



# Why So Many?



# Multiple John Singer Sargents?

**dallas:John\_Singer\_Sargent**

```
a foaf:Person;  
:dateOfBirth "1856" ;  
:dateOfDeath "1925" ;  
:name "John Singer Sargent" .
```

**ima:John\_Singer\_Sargent**

```
a foaf:Person;  
:dateOfBirth "1856" ;  
:dateOfDeath "1925" ;  
:name "John S. Sargent" .
```

# Multiple John Singer Sargents?

**dallas:John\_Singer\_Sargent**

```
a foaf:Person;  
:dateOfBirth "1856" ;  
:dateOfDeath "1925" ;  
:name "John Singer Sargent" .
```

) = ???

string\_match(

~~1 a a : John \_ Singer \_ Sargent~~

~~1 a a : John \_ Singer \_ Sargent~~

```
a foaf:Person;  
:dateOfBirth "1856" ;  
:dateOfDeath "1925" ;  
:name "John S. Sargent" .
```

# String Matching Problem

myMatchFunction(x, y)  
yourMatchFunction(x, y)

What does it mean that one is better than the other?



# Problem Definition

Given  $X$  and  $Y$  sets of strings

Find pairs  $(x, y)$   
such that both  $x$  and  $y$   
refer to the same real world entity

"John S. Sargent"

"John Singer Sargent"



# Problem Definition

Given  $X$  and  $Y$  sets of strings

Find pairs  $(x, y)$   
such that both  $x$  and  $y$   
refer to the same real world entity

We can use **precision** and **recall** to evaluate algorithms

# Problem Definition

Given  $X$  and  $Y$  sets of strings

Find pairs  $(x, y)$   
such that both  $x$  and  $y$   
refer to the same real world entity

fraction of pairs found that are correct  


We can use **precision** and **recall** to evaluate algorithms

 fraction of pairs found

# Why Strings Don't Match Perfectly?

typos "Joh" vs "John"

OCR errors "J0hn" vs "John"

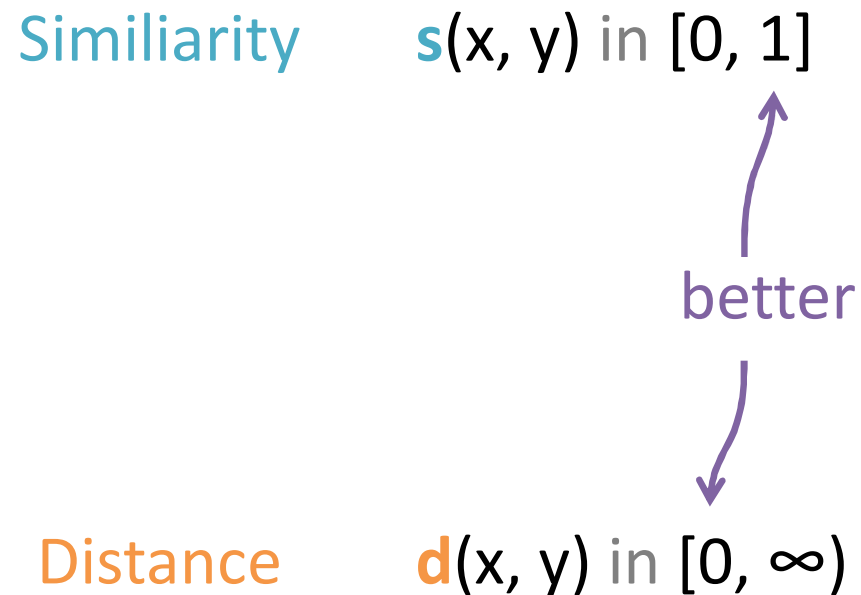
formatting conventions "03/17" vs "March 17"

abbreviations "J. S. Sargent" vs "John Singer Sargent"

nick names "John" vs "Jock"

word order "Sargent, John S." vs "John S. Sargent"

# Similarity Measure



# Types of Similarity Metrics

- Sequence based
- Set based
- Hybrid
- Phonetic

# Sequence Based Metrics

# Edit Distance

"J0n Singer Sargent"



"John S. Sargent"

insert character

delete character

substitute character


transpose character

...



# Edit Distance

"J0n Singer Sargent"



"John S. Sargent"

costs



insert character

$c_1$

delete character

$c_2$

substitute character

$c_3$

transpose character


$c_4$

...

$c_i$

# Edit Distance

"J0n Singer Sargent"



"John S. Sargent"

costs

insert character  $c_1$

delete character  $c_2$


substitute character  $c_3$

transpose character  $c_4$

...  $c_i$

# Levenshtein Distance

Edit distance:



insert character	1
delete character	1
substitute character	1

$\text{lev}(x, y)$  is the minimum cost to transform  $x$  to  $y$

Online calculator: <http://planetcalc.com/1721/>

# Computing Levenshtein Distance

$\text{lev}(x, y)$  is the minimum cost to transform  $x$  to  $y$

## Definitions

$$x = x_1 x_2 \dots x_n$$

$$y = y_1 y_2 \dots y_m$$

$$d(i, j) = \text{lev}(x_1 x_2 \dots \underline{x_i}, y_1 y_2 \dots \underline{y_j})$$

$$d(\emptyset, \emptyset) = \text{lev}(\text{""}, \text{""}) = 0$$

We want  $d(n, m)$

# Computing Levenshtein Distance

$$d(\emptyset, \emptyset) = 0$$

# Computing Levenshtein Distance

$$d(\emptyset, \emptyset) = 0$$

$x_1 x_2 \dots x_{i-1} x_i$  is a suffix of  $x$        $y_1 y_2 \dots y_{j-1} y_j$  is a suffix of  $y$

# Computing Levenshtein Distance

$$d(\emptyset, \emptyset) = 0$$

$x_1 x_2 \dots x_{i-1} x_i$  is a prefix of  $x$

$y_1 y_2 \dots y_{j-1} y_j$  is a prefix of  $y$

Case	Distance	Operation
$x_i = y_j$	$d(i-1, j-1)$	keep $x_i$

# Computing Levenshtein Distance

$$d(\emptyset, \emptyset) = 0$$

$x_1 x_2 \dots x_{i-1} x_i$  is a prefix of  $x$

$y_1 y_2 \dots y_{j-1} y_j$  is a prefix of  $y$

Case	Distance	Operation
$x_i = y_j$	$d(i-1, j-1)$	keep $x_i$
$x_i \neq y_j$		delete $x_i$ insert $y_j$ after $x_i$ replace $x_i$ with $y_j$



# Computing Levenshtein Distance

$$d(\emptyset, \emptyset) = 0$$

$x_1 x_2 \dots x_{i-1} x_i$  is a prefix of  $x$        $y_1 y_2 \dots y_{j-1} y_j$  is a prefix of  $y$

Case	Distance	Operation
$x_i = y_j$	$d(i-1, j-1)$	keep $x_i$
$x_i \neq y_j$	$d(i-1, j) + 1$	delete $x_i$ insert $y_j$ after $x_i$ replace $x_i$ with $y_j$

# Computing Levenshtein Distance

$$d(\emptyset, \emptyset) = 0$$

$x_1 x_2 \dots x_{i-1} x_i$  is a prefix of  $x$        $y_1 y_2 \dots y_{j-1} y_j$  is a prefix of  $y$

Case	Distance	Operation
$x_i = y_j$	$d(i-1, j-1)$	keep $x_i$
$x_i \neq y_j$	$d(i-1, j) + 1$ $d(i, j-1) + 1$	delete $x_i$ insert $y_j$ after $x_i$ replace $x_i$ with $y_j$

# Computing Levenshtein Distance

$$d(\emptyset, \emptyset) = 0$$

$x_1 x_2 \dots x_{i-1} x_i$  is a prefix of  $x$

$y_1 y_2 \dots y_{j-1} y_j$  is a prefix of  $y$

Case	Distance	Operation
$x_i = y_j$	$d(i-1, j-1)$	keep $x_i$
$x_i \neq y_j$	$d(i-1, j) + 1$	delete $x_i$
	$d(i, j-1) + 1$	insert $y_j$ after $x_i$
	$d(i-1, j-1) + 1$	replace $x_i$ with $y_j$

# Computing Levenshtein Distance

$$d(\emptyset, \emptyset) = 0$$

$x_1 x_2 \dots x_{i-1} x_i$  is a prefix of  $x$        $y_1 y_2 \dots y_{j-1} y_j$  is a prefix of  $y$

Case	Distance	Operation
$x_i = y_j$	$d(i-1, j-1)$	keep $x_i$
$x_i \neq y_j$	$d(i-1, j) + 1$	delete $x_i$
	$d(i, j-1) + 1$	insert $y_j$ after $x_i$
	$d(i-1, j-1) + 1$	replace $x_i$ with $y_j$

$$d(i, j) = \text{minimum} \uparrow$$

# Computing Levenshtein Distance Using Dynamic Programming

- $x = \text{dva}$ ,  $y = \text{dave}$

		y0	y1	y2	y3	y4
			<b>d</b>	<b>a</b>	<b>v</b>	<b>e</b>
x0		0	1	2	3	4
x1	<b>d</b>	1	0	1		
x2	<b>v</b>	2				
x3	<b>a</b>	3				

		y0	y1	y2	y3	y4
			d	a	v	e
x0		0	1	2	3	4
x1	d	1	0	1	2	3
x2	v	2	1	1	1	2
x3	a	3	2	1	2	2

$x = \text{d} - \text{v} \text{ a}$   
 $\quad \quad | \quad | \quad | \quad |$   
 $y = \text{d} \text{ a} \text{ v} \text{ e}$

substitute a with e  
 insert a (after d)

- Cost of dynamic programming is  $O(|x| |y|)$

# Levenshtein Distance Complexity

Dynamic programming algorithm

Time Complexity =  $O(N * M)$ , Space Complexity =  $O(N * M)$

## Polylogarithmic Approximation for Edit Distance and the Asymmetric Query Complexity

Alexandr Andoni, Robert Krauthgamer, Krzysztof Onak

We present a near-linear time algorithm that approximates the edit distance between two strings within a polylogarithmic factor; specifically, for strings of length  $n$  and every fixed  $\epsilon > 0$ , it can compute a  $(\log n)^{O(1/\epsilon)}$  approximation in  $n^{(1+\epsilon)}$  time.

<http://arxiv.org/abs/1005.4033>

# Levenshtein Distance Examples

lev(John Singer Sargent,  
John S. Sargent) =

# Levenshtein Distance Examples

`lev(John Singer Sargent,  
John S. Sargent)` = 5



# Levenshtein Distance Examples

lev(John Singer Sargent,  
John S. Sargent) = 5

lev(John Singer Sargent,  
Jane Klinger Sargent) =

# Levenshtein Distance Examples

lev(John Singer Sargent,  
John S. Sargent) = 5

lev(John Singer Sargent,  
Jane Klinger Sargent) = 5

# Levenshtein Distance Examples

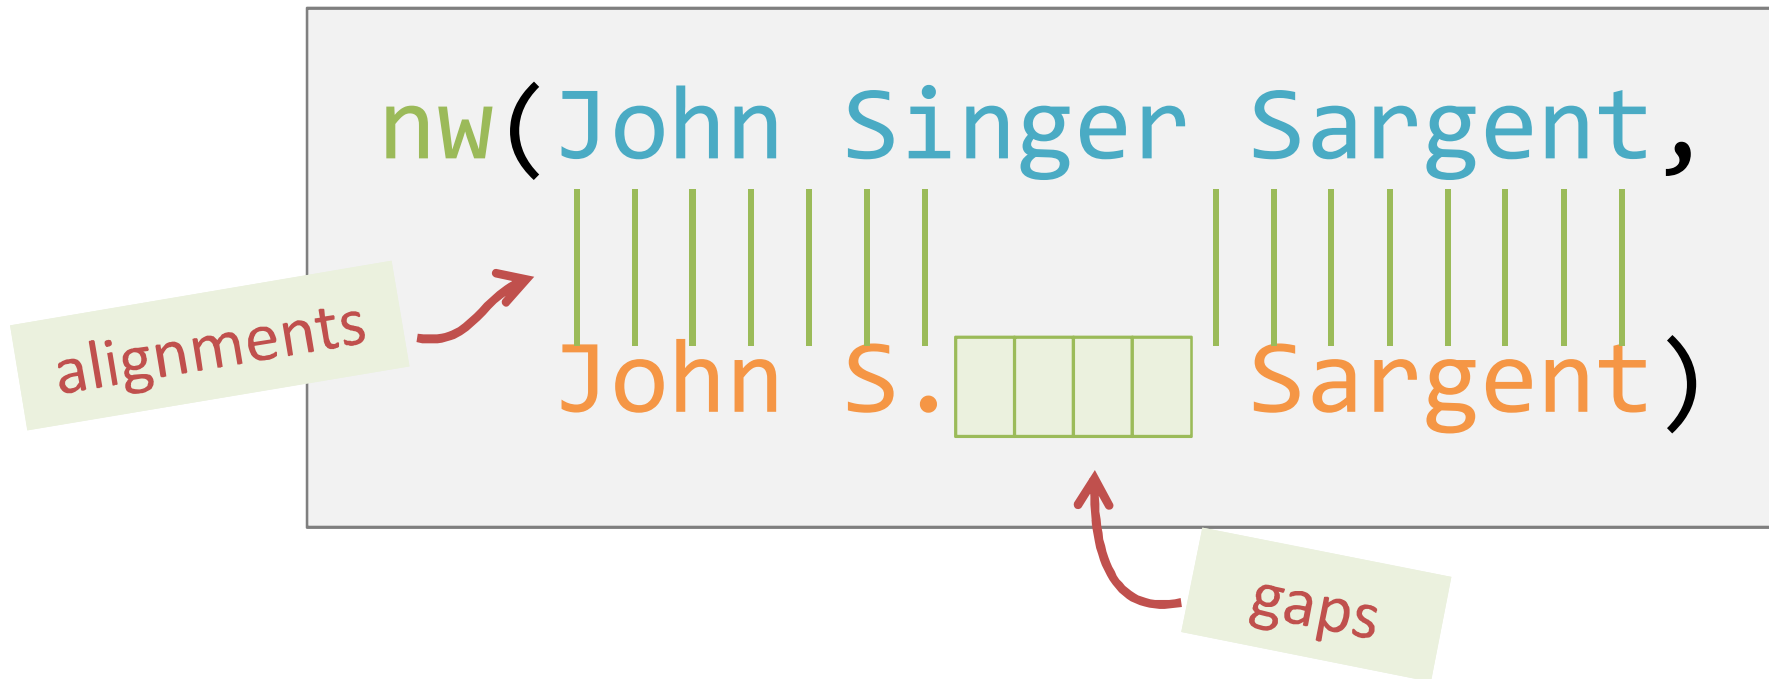
Too high a cost for deleting a  
sequence of characters


$$\text{lev}(\text{John Singer Sargent}, \text{John S. Sargent}) = 5$$

$$\text{lev}(\text{John Singer Sargent}, \text{Jane Klinger Sargent}) = 5$$

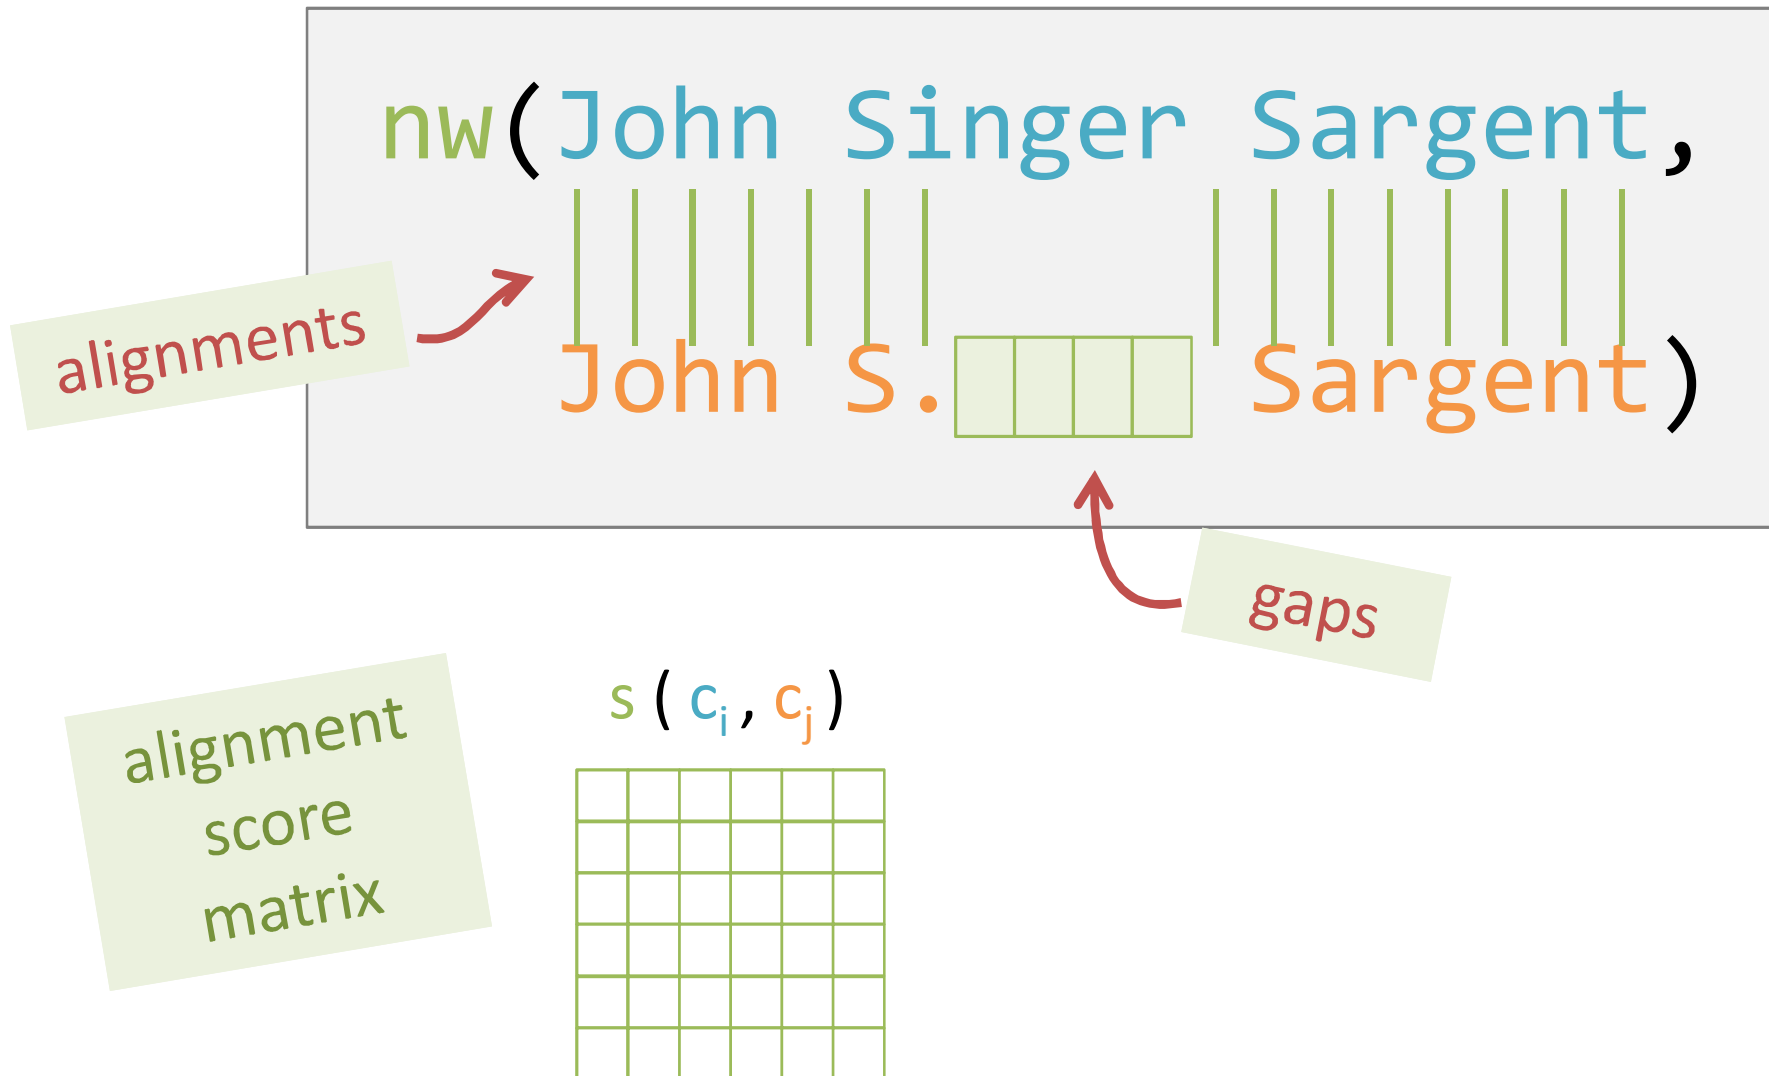
# Needleman-Wunch Measure

Generalization of `levenshtein(x, y)`



# Needleman-Wunch Measure

Generalization of `levenshtein(x, y)`



# Needleman-Wunch Measure

## Generalization of `levenshtein(x, y)`

nw(John Singer Sargent,

## alignments

John S. 

--	--	--	--

 Sargent)

gaps


## Alignment score matrix

$$s(c_i, c_j)$$
[illegible]

gap score

# Needleman-Wunch Measure

Generalization of `levenshtein(x, y)`



$$s(c_i, c_j) = \begin{cases} 2 & \text{if } c_i = c_j \\ -1 & \text{if } c_i \neq c_j \end{cases}$$

$$\text{gap-score} = -0.5$$

# Needleman-Wunch Measure

## Generalization of `levenshtein(x, y)`

$$2 * 14 + (-1) * 1 + (-0.5) * 4 = 25$$

$$s(c_i, c_j) = \begin{cases} 2 & \text{if } c_i = c_j \\ -1 & \text{if } c_i \neq c_j \end{cases}$$

gap-score = -0.5



# Comparison

	Levenshtein	Needleman-Wunch
Costs	1	matrix
Operations	insert/delete	gaps
Result	distance	similarity

OCR errors "J0hn" vs "John"

$\text{score}(\text{o}, \text{0}) = -0.2$

$\text{score}(\text{m}, \text{0}) = -1.0$

lower penalty



# Needleman-Wunch Example

nw(John Singer Sargent,  
John S. Sargent)

$$2 * 14 + (-1) * 1 + (-0.5) * 4 = 25$$

$$2 * 14 + (-1) * 3 + (-0.5) * 1 = 24.5$$

nw(John S inger Sargent,  
Jane Klinger Sargent)

# Needleman-Wunch Example

nw(John Singer Sargent,  
John S. Sargent)

$$2 * 14 + (-1) * 1 + (-0.5) * 4 = 25$$

$$2 * 14 + (-1) * 1 + (-0.5) * 8 = 23$$

nw(John Stanislaus Sargent,  
John S. Sargent)

# Needleman-Wunch Example

nw(John Singer Sargent,  
John S. Sargent)

$$+ (-1) * 1 + (-0.5) * 4 = 25$$

Longer gaps are  
penalized more


Bad for names

$$-2 - 14 + (-1) * 1 + (-0.5) * 8 = 23$$

nw(John Stanislaus Sargent,  
John S. Sargent)


# Affine Gap Measure

Generalization of `needleman-wunch(x, y)`



# Affine Gap Measure

Generalization of `needleman-wunch(x, y)`



continue gap, cost  $c_r$

open gap, cost  $c_o$

total gap cost:  $c_o + (|gap| - 1) * c_r$

# Smith-Waterman

## Global alignment

fully align sequences,  
opening gaps as needed

F	T	F	T	A	L	I	L	L	A	V	A	V
F	-	-	T	A	L	-	L	L	A	-	A	V

## Needleman-Wunch

## Local alignment

find best subsequences  
to align

F	T	F	T	A	L	I	L	L	A	V	A	V
-	-	F	T	A	L	-	L	L	A	V	-	-

Smith-Waterman: local alignment version of Needleman-Wunch

# Smith-Waterman Example

```
match(John Sargent, american painter,  
      American artist John S. Sargent)
```

```
nw(  
    John S    argent, american painter,  
    American artist John S. Sargent    )
```

Needleman-Wunch: significant gap penalty



# Smith-Waterman Example

```
match(John Sargent, american painter,  
      American artist John S. Sargent)
```

```
nw(
  American artist John S. Sargent, american painter,
)
```

## Needleman-Wunch: significant gap penalty

# Smith-Waterman Example

```
match(John Sargent, american painter,  
      American artist John S. Sargent)
```

```
nw(
  American artist John S. Sargent, american painter,
)
```

## Needleman-Wunch: significant gap penalty

```
sw(John Sargent, american painter,  
    American artist John S. Sargent)
```

## Smith-Waterman: identifies similar subsequences

# Jaro Similarity Measure

- Get points for having characters in common
  - but only if they are “close by”
- Get points for common characters in the same order
  - lose points for transpositions

# Jaro Similarity Measure $\text{jaro}(x, y)$

$$\text{max-distance} = \frac{\max(|x|, |y|)}{2} - 1$$

$$x_i \text{ matches } y_j \text{ if } \begin{cases} x_i = y_j \\ |i - j| \leq \text{max-distance} \end{cases}$$

$m$  = number of matching characters

$t$  = number of transpositions  
(of matching characters)

# Jaro Similarity Measure

$$\text{max-distance} = \frac{\max(|x|, |y|)}{2} - 1$$

$m$  = number of matching characters

$t$  = number of transpositions

$$\text{jaro}(x, y) = \begin{cases} 0 & \text{if } m = 0 \\ \frac{1}{3} \left( \frac{m}{|x|} + \frac{m}{|y|} + \frac{m - t}{m} \right) & \text{otherwise} \end{cases}$$

# Jaro Example

lev(DIXON, DICKSONX) = 4      (4/8 = 0.5)

jaro(DIXON, DICKSONX) = ???

# Jaro Example

	D	I	X	O	N
D	1	0	0	0	0
I	0	1	0	0	0
C	0	0	0	0	0
K	0	0	0	0	0
S	0	0	0	0	0
O	0	0	0	1	0
N	0	0	0	0	1
X	0	0	1	0	0

# Jaro Example

$$|x| = 5$$

$$|y| = 8$$

$$\begin{aligned} \text{max-distance} &= (8/2) - 1 \\ &= 3 \end{aligned}$$

$$m = 4$$

$$t = 0$$

	D	I	X	O	N
D	1	0	0	0	0
I	0	1	0	0	0
C	0	0	0	0	0
K	0	0	0	0	0
S	0	0	0	0	0
O	0	0	0	1	0
N	0	0	0	0	1
X	0	0	0	0	0



# Jaro Example

	D	I	X	O	N
D	1	0	0	0	0
I	0	1	0	0	0
C	0	0	0	0	0
K	0	0	0	0	0
S	0	0	0	0	0
O	0	0	0	1	0
N	0	0	0	0	1
X	0	0	0	0	0

$$|x| = 5$$

$$|y| = 8$$

$$m = 4$$

$$t = 0$$

$$\frac{1}{3} \left( \frac{m}{|x|} + \frac{m}{|y|} + \frac{m - t}{m} \right)$$

$$\frac{1}{3} \left( \frac{4}{5} + \frac{4}{8} + \frac{4 - 0}{4} \right)$$

$$= 0.767$$

# Jaro Example

$$\text{lev}(\text{DIXON}, \text{DICKSONX}) = 4 \quad (4/8 = 0.5)$$

$$\text{jaro}(\text{DIXON}, \text{DICKSONX}) = 0.767$$

# Jaro-Winkler Measure

Give a bonus if there is a common prefix

```
jwProximity(x,y,boostThreshold,prefixSize)
= jaro(x,y) <= boostThreshold
? jaro (x,y)
: jaro (x,y)
  + 0.1 * prefixMatch(x,y,prefixSize)
    * (1.0 - jaro(x,y))
```

```
prefixMatch(x,y,prefixSize) =
  min(prefixSize, common-prefix(x,y))
```

```
boostThreshold = 0.7
prefixSize = 4
```

# Jaro-Winkler Measure

```
jwProximity(x,y,boostThreshold,prefixSize)
  = jaro(x,y) <= boostThreshold
  ? jaro (x,y)
  : jaro (x,y)
    + 0.1 * prefixMatch(x,y,prefixSize)
      * (1.0 - jaro(x,y))
```

```
boostThreshold = 0.7
prefixSize = 4
```

```
jaro(DIXON, DICKSONX) = 0.767
```

```
jwProximity(DIXON, DICKSONX) = 0.767 + 0.1*2*(1 - 0.767)
                                = 0.767 + 0.2*0.233
                                = 0.767 + 0.0466
                                = 0.8136
```

# Set-Based Metrics

# Set-Based Metrics

Generate set of tokens from the strings

Measure similarity between the sets of tokens

# Tokenizing a String

Words

# Tokenizing a String

Words

$q$ -grams: substrings of length  $q$

“david smith” 3-grams  
{##d, #da, dav, avi, ..., h##}



# Jaccard Measure

$B_x = \text{tokens}(x)$

$B_y = \text{tokens}(y)$

$$\text{jaccard}(x, y) = \frac{|B_x \cap B_y|}{|B_x \cup B_y|}$$

$\text{jaccard}(\text{dave}, \text{dav})$

$B_x = \{\#d, da, av, ve, e\# \}$

$B_y = \{\#d, da, av, v\# \}$

$\text{jaccard}(x, y) = 3/6$

# TF/IDF Measure

TF = term frequency

IDF = inverse document frequency

$x$  = Apple Corporation, CA

$y$  = IBM Corporation, CA

$z$  = Apple Corp

...

blah blah Corporation

$\text{lev}(x, y)$  =  $\text{lev}(x, z)$

>

<

???

# TF/IDF Measure

TF = term frequency

IDF = inverse document frequency

$x$  = Apple Corporation, CA

$y$  = IBM Corporation, CA

$z$  = Apple Corp

...

blah blah Corporation

$\text{lev}(x, y) > \text{lev}(x, z)$

... but intuitively  $(x, z)$  is a better match

# TF/IDF Measure

TF = term frequency

IDF = inverse document frequency

x = Apple Corporation, CA  
y = IBM Corporation, CA  
z = Apple Corp  
...  
blah blah Corporation

frequent term  
should not carry  
as much weight

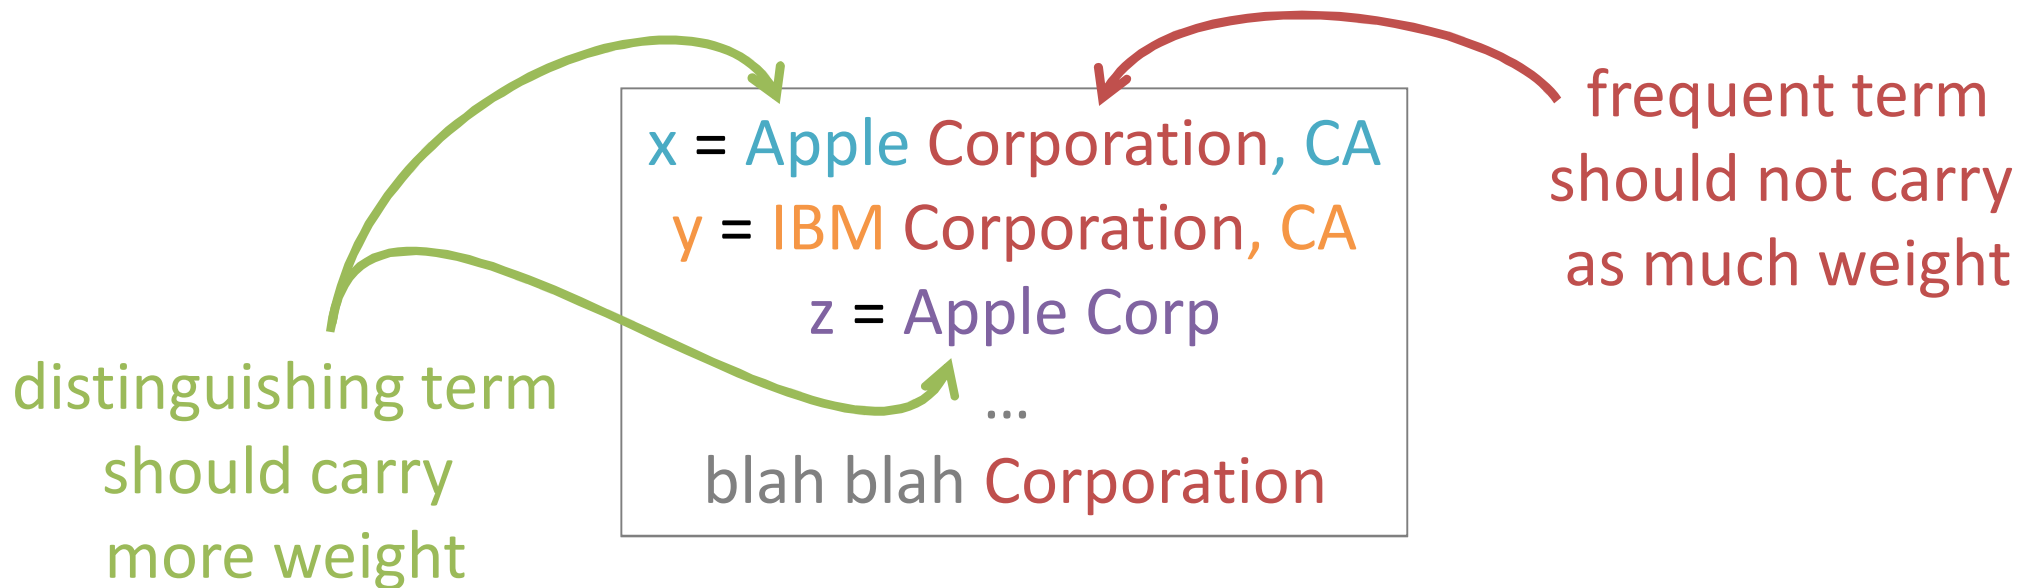
$$\text{lev}(x, y) > \text{lev}(x, z)$$

... but intuitively (x, z) is a better match

# TF/IDF Measure

TF = term frequency

IDF = inverse document frequency



$$\text{lev}(x, y) > \text{lev}(x, z)$$

... but intuitively  $(x, z)$  is a better match

# Term Frequencies and Inverse Document Frequencies

- Assume  $x$  and  $y$  are taken from a collection of strings
- Each string is converted into a bag of terms called a document
- term frequency  $tf(t,d) =$ 
  - number of times term  $t$  appears in document  $d$
- inverse document frequency  $idf(t) =$ 
  - $N / N_d$ , number of documents in collection divided by number of documents that contain  $t$
  - note: in practice,  $idf(t)$  is often defined as  $\log(N / N_d)$

# Example

$$x = aab \Rightarrow B_x = \{a, a, b\}$$

$$y = ac \Rightarrow B_y = \{a, c\}$$

$$z = a \Rightarrow B_z = \{a\}$$

$$\text{tf}(a, x) = 2 \qquad \text{idf}(a) = 3/3 = 1$$

$$\text{tf}(b, x) = 1 \qquad \text{idf}(b) = 3/1 = 3$$

$$\dots \qquad \text{idf}(c) = 3/1 = 3$$

$$\text{tf}(c, z) = 0$$

# Feature Vectors

- Each document  $d$  is converted into a feature vector  $\mathbf{v}_d$
- $\mathbf{v}_d$  has a feature  $v_d(t)$  for each term  $t$ 
  - value of  $v_d(t)$  is a function of TF and IDF scores
  - here we assume  $v_d(t) = \text{tf}(t,d) * \text{idf}(t)$

$x = \text{aab} \Rightarrow B_x = \{a, a, b\}$

$y = \text{ac} \Rightarrow B_y = \{a, c\}$

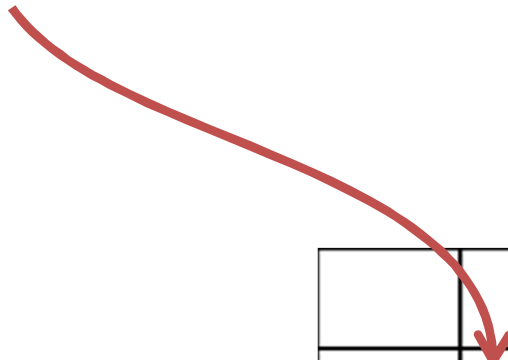
$z = a \Rightarrow B_z = \{a\}$

$\text{tf}(a, x) = 2$        $\text{idf}(a) = 3/3 = 1$

$\text{tf}(b, x) = 1$        $\text{idf}(b) = 3/1 = 3$

...       $\text{idf}(c) = 3/1 = 3$

$\text{tf}(c, z) = 0$



	<b>a</b>	<b>b</b>	<b>c</b>
<b>v<sub>x</sub></b>	2	3	0
<b>v<sub>y</sub></b>	3	0	3
<b>v<sub>z</sub></b>	3	0	0



# TF/IDF Similarity Score

- Let  $p$  and  $q$  be two strings, and  $T$  be the set of all terms in the collection
- Feature vectors  $\mathbf{v}_p$  and  $\mathbf{v}_q$  are vectors in the  $|T|$ -dimensional space where each dimension corresponds to a term
- TF/IDF score of  $p$  and  $q$  is the cosine of the angle between  $\mathbf{v}_p$  and  $\mathbf{v}_q$ 
  - $s(p,q) = \sum_{t \in T} v_p(t) * v_q(t) / [\sqrt{\sum_{t \in T} v_p(t)^2} * \sqrt{\sum_{t \in T} v_q(t)^2}]$

# TF/IDF Similarity Score


- Score is high if strings share many frequent terms
  - terms with high TF scores
- Unless these terms are common in other strings
  - i.e., they have low IDF scores
- Dampening TF and IDF as commonly done in practice
  - use  $v_d(t) = \log(\text{tf}(t,d) + 1) * \log(\text{idf}(t))$  instead of  $v_d(t) = \text{tf}(t,d) * \text{idf}(t)$
- Normalizing feature vectors

- $$v_d(t) = v_d(t) / \sqrt{\sum_{\{t \in T\}} v_d(t)^2}$$

# Hybrid Similarity Measures

# Hybrid Measures

Do the set-based thing  
but  
use a similarity metric for each element of the set




# Generalized Jaccard Measure

- Jaccard measure
  - considers overlapping tokens in both x and y
  - a token from x and a token from y must be identical to be included in the set of overlapping tokens
  - this can be too restrictive in certain cases
- Example:
  - matching taxonomic nodes that describe companies
  - “Energy & Transportation” vs. “Transportation, Energy, & Gas”
  - in theory Jaccard is well suited here, in practice Jaccard may not work well if tokens are commonly misspelled
    - e.g., energy vs. eneryg
  - generalized Jaccard measure can help such cases

# Generalized Jaccard Measure

- Let  $B_x = \{x_1, \dots, x_n\}$ ,  $B_y = \{y_1, \dots, y_m\}$
- Step 1: find token pairs that will be in the “softened” overlap set
  - apply a similarity measure  $s$  to compute sim score for each pair  $(x_i, y_j)$
  - keep only those score  $>$  a given threshold  $\alpha$ , this forms a bipartite graph  $G$
  - find the maximum-weight matching  $M$  in  $G$
- Step 2: return normalized weight of  $M$  as generalized Jaccard score
  - $GJ(x,y) = \sum_{(x_i,y_j) \text{ in } M} s(x_i,y_j) / (|B_x| + |B_y| - |M|)$

# Generalized Jaccard Example



$$\alpha = 0.5$$

- Generalized Jaccard score:  $(0.7 + 0.9)/(3 + 2 - 2) = 0.53$

# The Soft TF/IDF Measure


- Similar to generalized Jaccard measure, except that it uses TF/IDF measure as the “higher-level” sim measure
  - e.g., “Apple Corporation, CA”, “IBM Corporation, CA”, and “Aple Corp”, with Apple being misspelled
- Step 1: compute  $\text{close}(x,y,k)$ :  
set of all terms  $t \in B_x$  that have at least one close term  $u \in B_y$ , i.e.,  $s'(t,u) \geq k$ 
  - $s'$  is a basic sim measure (e.g., Jaro-Winkler),  $k$  prespecified
- Step 2: compute  $s(x,y)$  as in traditional TF/IDF score, but weighing each TF/IDF component using  $s'$ 
  - $s(x,y) = \sum_{t \in \text{close}(x,y,k)} v_x(t) * v_y(u^*) * s'(t,u^*)$
  - $u^* \in B_y$  maximizes  $s'(t,u) \forall u \in B_y$



# Soft TF/IDF Example

$x = abcd$

$y = aa'b'cd'$



$\text{close}(x, y, 0.75) = \{a, b, c\}$

$$s(x, y) = v_x(a) \cdot v_y(a) \cdot 1 + v_x(b) \cdot v_y(b') \cdot 0.8 + v_x(c) \cdot v_y(c) \cdot 1$$

(a)

(b)

(c)

# Monge-Elkan Measure

- Break strings  $x$  and  $y$  into multiple substrings
  - $x = A_1 \dots A_n$ ,  $y = B_1 \dots B_m$
- Compute
  - $s(x,y) = 1/n * \sum_{i=1}^n \max_{j=1}^m s'(A_i, B_j)$
  - $s'$  is a secondary sim measure, such as Jaro-Winkler
  - Intuitively, we ignore the order of the matching of substrings and only consider the best match for substrings of  $x$  in  $y$

# Monge-Elkan Measure

$$s(x,y) = 1/n * \sum_{i=1}^n \max_{j=1}^m s'(A_i, B_j)$$

$$x = A_1 A_2$$

$$y = B_1 B_2 B_3$$

$$\frac{\begin{aligned} &\max( s'(A_1, B_1), s'(A_1, B_2), s'(A_1, B_3) ) \\ &+ \\ &\max( s'(A_2, B_1), s'(A_2, B_2), s'(A_2, B_3) ) \end{aligned}}{2}$$

$s'$  could be any metric, e.g., levenshtein

# Monge-Elkan Measure

$$s(x,y) = 1/n * \sum_{i=1}^n \max_{j=1}^m s'(A_i, B_j)$$

x = Comput. Sci. and Eng. Dept., University of California, San Diego

y = Department of Computer Science, Univ. of Calif., San Diego

what  $s'$  should we use?

levenshtein

needleman-wunch

affine-gap

smith-waterman

jaro

jaro-winkler

# Phonetic Similarity Measures

# Phonetic Similarity Measures

- Match strings based on their sound, instead of appearances
- Very effective in matching names, which often appear in different ways that sound the same
  - e.g., Meyer, Meier, and Mire; Smith, Smithe, and Smythe
- Soundex is most commonly used

# The Soundex Measure

- Used primarily to match surnames
  - maps a surname  $x$  into a 4-letter code
  - two surnames are judged similar if share the same code
- Algorithm to map  $x$  into a code:
  - Step 1: keep the first letter of  $x$ , subsequent steps are performed on the rest of  $x$
  - Step 2: remove all occurrences of W and H. Replace the remaining letters with digits as follows:
    - ❖ replace B, F, P, V with 1, C, G, J, K, Q, S, X, Z with 2, D, T with 3, L with 4, M, N with 5, R with 6
  - Step 3: replace sequence of identical digits by the digit itself
  - Step 4: Drop all non-digit letters, return the first four letters as the soundex code

# The Soundex Measure

- Example: x = Ashcraft
  - after Step 2: A226a13, after Step 3: A26a13, Step 4 converts this into A2613, then returns A261
  - Soundex code is padded with 0 if there is not enough digits
- Example: Robert and Rupert map into R163
- Soundex fails to map Gough and Goff, and Jawornicki and Yavornitzky
  - designed primarily for Caucasian names, but found to work well for names of many different origins
  - does not work well for names of East Asian origins
    - ❖ which uses vowels to discriminate, Soundex ignores vowels



# Other Readings

- [http://en.wikipedia.org/wiki/String\\_metric](http://en.wikipedia.org/wiki/String_metric)
  - [http://en.wikipedia.org/wiki/Approximate\\_string\\_matching](http://en.wikipedia.org/wiki/Approximate_string_matching)
  - [http://en.wikipedia.org/wiki/Edit\\_distance](http://en.wikipedia.org/wiki/Edit_distance)
  - [http://en.wikipedia.org/wiki/Levenshtein\\_distance](http://en.wikipedia.org/wiki/Levenshtein_distance)
  - [http://en.wikipedia.org/wiki/Jaro-Winkler\\_distance](http://en.wikipedia.org/wiki/Jaro-Winkler_distance)
  - [http://en.wikipedia.org/wiki/Smith-Waterman\\_algorithm](http://en.wikipedia.org/wiki/Smith-Waterman_algorithm)
  - [http://en.wikipedia.org/wiki/Jaccard\\_index](http://en.wikipedia.org/wiki/Jaccard_index)
- 
- <http://alias-i.com/lingpipe/demos/tutorial/stringCompare/read-me.html>
  - <http://www.gettingcirrius.com/2011/01/calculating-similarity-part-2-jaccard.html>
  - [http://en.wikipedia.org/wiki/Sequence\\_alignment#Pairwise\\_alignment](http://en.wikipedia.org/wiki/Sequence_alignment#Pairwise_alignment)

# Software

- <http://code.google.com/p/java-similarities/source/browse/trunk/simmetrics/src/main/java/uk/ac/shef/wit/simmetrics/>
- <http://sourceforge.net/projects/secondstring/>
- <http://planetcalc.com/1721/> (Levenshtein calculator)

**Source path:** [svn/](#) [trunk/](#) [simmetrics/](#) [src/](#) [main/](#) [java/](#) [uk/](#) [ac/](#) [shef/](#) [wit/](#) [simmetrics](#)

Directories	Filename
▼simmetrics	<a href="#">AbstractStringMetric.java</a>
arbitrators	<a href="#">BlockDistance.java</a>
basiccontainers	<a href="#">ChapmanLengthDeviation.java</a>
math	<a href="#">ChapmanMatchingSoundex.java</a>
metrichandlers	<a href="#">ChapmanMeanLength.java</a>
▼similaritymetrics	<a href="#">ChapmanOrderedNameCompoundSimilarity.java</a>
costfunctions	<a href="#">CosineSimilarity.java</a>
task	<a href="#">DiceSimilarity.java</a>
tokenisers	<a href="#">EuclideanDistance.java</a>
utils	<a href="#">InterfaceStringMetric.java</a>
wordhandlers	<a href="#">JaccardSimilarity.java</a>
	<a href="#">Jaro.java</a>
	<a href="#">JaroWinkler.java</a>
	<a href="#">Levenshtein.java</a>
	<a href="#">MatchingCoefficient.java</a>
	<a href="#">MongeElkan.java</a>
	<a href="#">NeedlemanWunch.java</a>
	<a href="#">OverlapCoefficient.java</a>
	<a href="#">QGramsDistance.java</a>