

Lecture 16: March 21, 2015
cs 573: Probabilistic Reasoning
Professor Nevatia
Spring 2015

Review

- HW #5 due Mar 30;
- Exam 1, please return with or without comments
- Exam 2, April 29, class period; NOT cumulative
- Previous Lecture
 - Variational approximation
- Today's objective
 - Review and complete Variational Approximation discussion
 - Sampling approach

Variational Approach

- Treats inference problem as an optimization problem
- Approximate the actual distribution, say P , with a simpler distribution, say Q
 - e.g. fit a 1-D Gaussian to a 1-D arbitrary distribution
 - Fit a multi-variate Gaussian but assume co-variances are nil
 - Assume Q is a product of individual variable distributions
 - Need to select a tractable Q
- How to measure difference between P and Q ?
- How to choose parameters of Q to minimize the difference between P and Q ?

Finding the Minimum

- We choose to minimize $D(Q||P_\Phi) = -H_Q(X) - E_Q[\ln P_\Phi(X)]$

Equivalent to maximizing “free energy” functional

$$F[\tilde{P}_\Phi, Q] = H_Q(X) + \sum_{\phi \in \Phi} E_Q[\ln \phi(U_\phi)] \text{ Find } Q \text{ that minimizes}$$

- Choose $Q(X) = \prod_i Q(X_i)$; also $\sum_i Q(X_i) = 1$
- Take derivatives of the Lagrangian with respect to $Q(x_i)$ and set $= 0$
- Leads to a “Mean-Field” approximation

$$Q(x_i) = \frac{1}{Z_i} \exp \{ E_{X_{-i} \sim Q} [\ln P_\Phi(x_i | X_{-i})] \}$$

- After further simplification:


$$Q(x_i) = \frac{1}{Z_i} \exp \left\{ \sum_{\phi: X_i \in \text{Scope}[\phi]} E_{(U_\phi - \{X_i\}) \sim Q} [\ln \phi(U_\phi, x_i)] \right\}.$$

where Z_i is a normalizing constant.

Explaining the Notation

- Start with a clique of just two variables (X_1, X_2) , given $\phi(X_1, X_2)$
- For X_1 : $Q(X_1=x_1^i) = 1/z_1 \exp (\sum_{x_2} Q(x_2) \log \phi(x_1^i, x_2))$
 - Sum is over all possible values of X_2
 - Note above updates the distribution $Q(X_1)$
 - Similar expression for $Q(X_2)$ in terms of $Q(X_1)$; iterate
- Suppose we have a clique (X_1, X_2, X_3)
For X_1 : $Q(X_1=x_1^i) = 1/z_1 \exp (\sum_{x_2, x_3} Q(x_2, x_3) \log \phi(x_1^i, x_2, x_3))$
 $= 1/z_1 \exp (\sum_{x_2, x_3} Q(x_2)Q(x_3) \log \phi(x_1^i, x_2, x_3))$
- If X_1 is also in another clique, say (X_1, X_4) , add terms corresponding to $\sum_{x_4} Q(x_4) \log \phi(x_1, x_4)$ to the above summation

GRID MRF



Grid Network Example

- Q for each node is a product of four factors in which it appears (sum in the log notation)

$$Q(a_{i,j}) = \frac{1}{Z_{i,j}} \exp \left\{ \begin{array}{l} \sum_{a_{i-1,j}} Q(a_{i-1,j}) \ln(\phi(a_{i-1,j}, a_{i,j})) + \\ \sum_{a_{i,j-1}} Q(a_{i,j-1}) \ln(\phi(a_{i,j-1}, a_{i,j})) + \\ \sum_{a_{i+1,j}} Q(a_{i+1,j}) \ln(\phi(a_{i,j}, a_{i+1,j})) + \\ \sum_{a_{i,j+1}} Q(a_{i,j+1}) \ln(\phi(a_{i,j}, a_{i,j+1})) \end{array} \right\}.$$

- The term in the exponent is the weighted “mean” of values of neighboring nodes; weight is function of affinity
- Iterative algorithm (11.7), see next slide
- Note: energy always decreases so algorithm is guaranteed to converge, but only to a local minimum

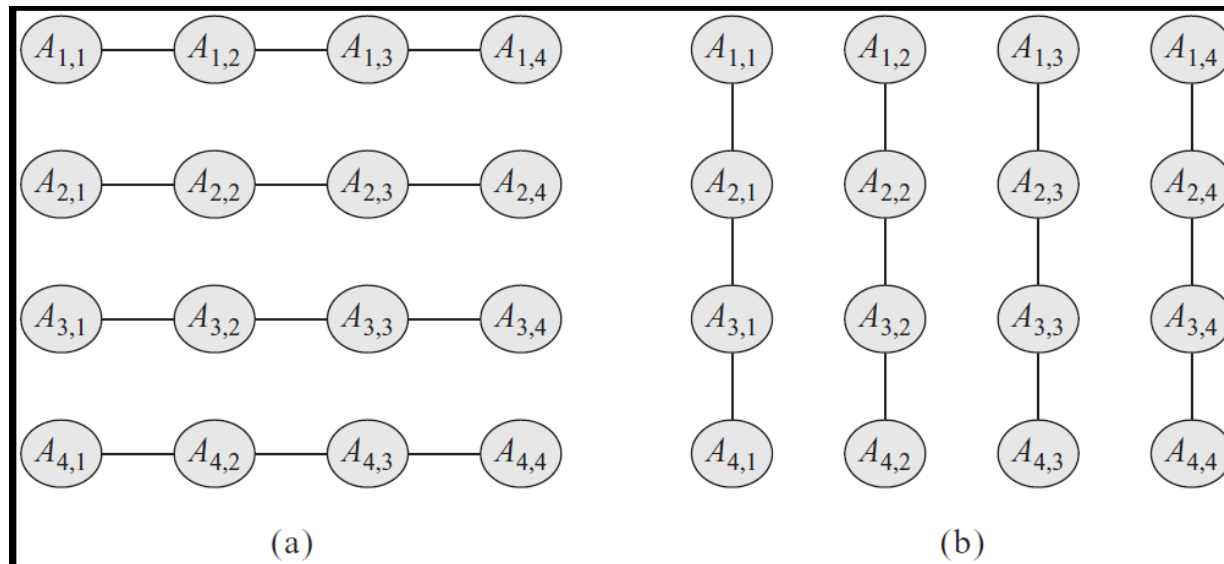
Algorithm 11.7

Algorithm 11.7 The Mean-Field approximation algorithm

```
Procedure Mean-Field (  
     $\Phi$ ,    // factors that define  $P_\Phi$   
     $Q_0$     // Initial choice of  $Q$   
)  
1    $Q \leftarrow Q_0$   
2    $Unprocessed \leftarrow \mathcal{X}$   
3   while  $Unprocessed \neq \emptyset$   
4       Choose  $X_i$  from  $Unprocessed$   
5        $Q_{old}(X_i) \leftarrow Q(X_i)$   
6       for  $x_i \in Val(X_i)$  do  
7            $Q(x_i) \leftarrow \exp \left\{ \sum_{\phi: X_i \in Scope[\phi]} E_{(U_\phi - \{X_i\}) \sim Q} [\ln \phi[U_\phi, x_i]] \right\}$   
8       Normalize  $Q(X_i)$  to sum to one  
9       if  $Q_{old}(X_i) \neq Q(X_i)$  then  
10           $Unprocessed \leftarrow Unprocessed \cup (\cup_{\phi: X_i \in Scope[\phi]} Scope[\phi])$   
11           $Unprocessed \leftarrow Unprocessed - \{X_i\}$   
12  return  $Q$ 
```


Structured Approximations

- Approximating distribution can be more complex (but still allow for efficient, exact inferences); see example below



- It is not difficult to derive conditions for maximizing the energy functional for any Q represented as a Gibbs distribution (eq. 11.61)
- Updates over the factors containing cliques in Q
 - May require inferences on the Q distribution after each step
 - In our earlier example, inference was easy as each node was not connected to the others
- We skip the details of this process

Structured Approximations vs LBP

- Structured Approximations
 - Iterative procedure with guarantee of convergence
 - Convergence may be to a local minimum
 - Trade-off between expressive power of the approximating function and computation time
- LBP
 - Easy to understand and implement
 - No guarantee of convergence
 - Seems to work “well” for many practical problems
- In general, empirical evaluation can be difficult as we may not be able to find the optimal, even for comparisons, due to inherent computational complexity.

Particle Based Approximate Inference

- In this approach, we simulate and collect specific instantiations of the variables, called *particles*, according to the network distributions, and use them to approximate the overall distribution.
- If we have “enough” particles, we can get a “good” approximation of the distribution function.
- Task will be to estimate some distribution $P(\mathbf{Y} = \mathbf{y})$ (note \mathbf{Y} can be a set of variables)
 - More generally, compute the expectation (expected value) of some function $f(\mathbf{X})$.
- Notation: $\xi \langle \mathbf{Y} \rangle$ is the assignment in ξ to the variables \mathbf{Y}
 - Note: ξ is pronounced as “ksi” but we often drop “k”
 - Choose $f(\xi) = \mathbf{1} \{ \xi \langle \mathbf{Y} \rangle = \mathbf{y} \}$ to compute $P(\mathbf{Y} = \mathbf{y})$

Sampling a single discrete random variable

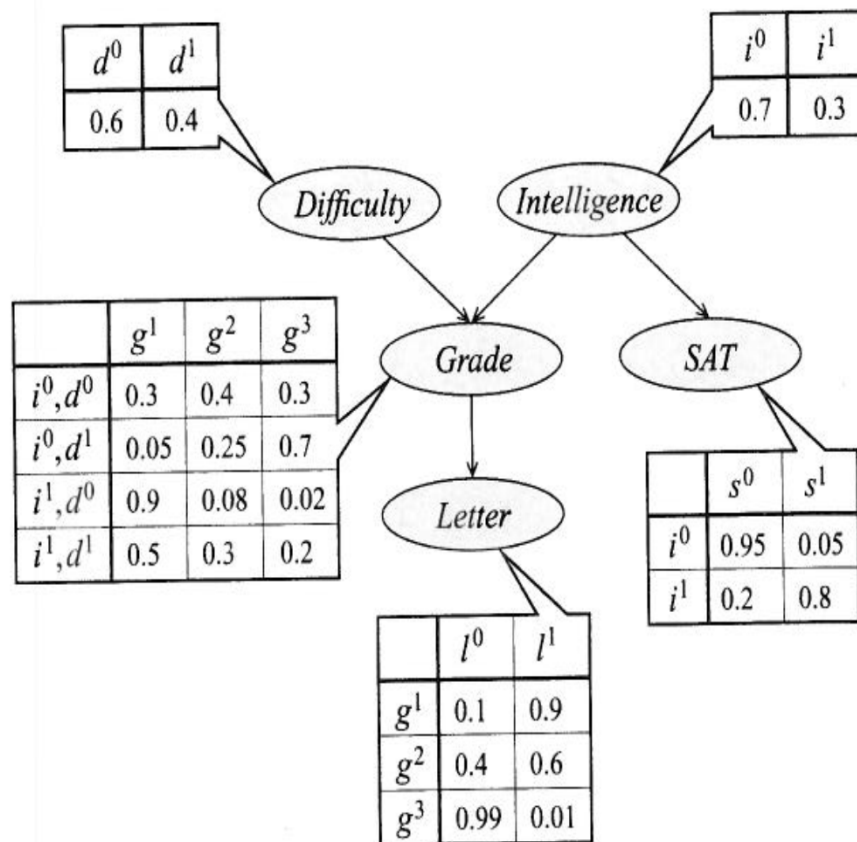
- How to generate random samples of x given $\mathbf{P}(x)$.
 - Consider a fair coin toss, for example
 - Number of H and T results s/b equal over a very long sequence
 - However, alternating sequence is of low probability; many unbalanced sequences should occur with varying probabilities
 - Coin has no memory; how to simulate its behavior?
- One approach is to sample from a uniform distribution random number generator in range of $[0, 1]$ and map to desired $P(X)$
- How to sample from the uniform distribution?
 - Observe a physical phenomenon satisfying this law
 - Pseudo-random number generator
 - Deep mathematical topic; Example: $X_{n+1} = (aX_n + c) \bmod m$
 - Choice of parameters determines “quality”
 - In C/C++ , `srand (time)` to initialize, `rand(1)` to get first number etc.
 - Mersenne Twister (high quality, fast random number generator)

Mapping to Specific Distributions

- Consider X to represent results of a coin toss
 - Declare Head if $r < 0.5$, Tail otherwise for fair coin
 - Let $P(\text{head}) = 0.7$, then declare Head if $r < 0.7$
- Consider X to be sum of pair of dice, all probabilities are not necessarily equal
 - Make ranges proportional to probabilities of the 12 outcomes
- Continuous distributions: let $F(x)$ be the cumulative distr function
- Let a sample from $[0,1]$ have value z , then:
- Choose smallest x for which $F(x) > z$, *i.e.* compute $F^{-1}(z)$
 - This provides the desired sample (Proof omitted)
- Computing $F^{-1}(z)$ may be difficult for some functions, such as a Gaussian; numerical methods may need to be used
 - More specialized methods exist for Gaussian functions
 - *e.g.* Add multiple uniform random numbers

Forward Sampling

- Generate random samples $\xi[1], \xi[2], \dots, \xi[M]$ from $P(X)$.
- Consider a BN, $P_B(X)$; example below



First sample D, let $D = d^1$

Sample I, let $I = i^0$

Sample G, from $P(G|i^0, d^1)$
and so on

See Alg 12.1

Forward Sampling Algorithm

Algorithm 12.1 Forward Sampling in a Bayesian network

Procedure Forward-Sample (
 \mathcal{B} // Bayesian network over \mathcal{X}
)

- 1 Let X_1, \dots, X_n be a topological ordering of \mathcal{X}
- 2 **for** $i = 1, \dots, n$
- 3 $u_i \leftarrow \mathbf{x} \langle \text{Pa}_{X_i} \rangle$ // Assignment to Pa_{X_i} in x_1, \dots, x_{i-1}
- 4 Sample x_i from $P(X_i \mid u_i)$
- 5 **return** (x_1, \dots, x_n)

Expectation from Samples

- From the set of particles, $D = \{\xi[1], \xi[2], \dots, \xi[M]\}$, we can compute expectation of any function of samples:

$$\hat{E}_D(f) = \frac{1}{M} \sum_{m=1}^M f(\xi[m]).$$

- If we want to compute $P(\mathbf{y})$, f is the indicator function equal to 1 when assignments to \mathbf{Y} in $\xi[M] = \mathbf{y}$; alternate notation is $\mathbf{y}[m]$ or $\xi[M] <\mathbf{y}>$

$$\hat{P}_D(\mathbf{y}) = \frac{1}{M} \sum_{m=1}^M \mathbf{I}\{\mathbf{y}[m] = \mathbf{y}\},$$

- $P(G = g^1, i^0) = (1/M) * \#(D=?, I=i^0, G=g^1, L=?, S=?)$
- $P(G = g^1 | i^0) = \#(D=?, I=i^0, G=g^1, L=?, S=?)/$
 $\#(D=?, I=i^0, G=?, L=?, S=?)$

Error Analysis

- Absolute error ($P(\mathbf{y}) \pm \epsilon$) or relative error ($P(\mathbf{y})$ is between estimate $\rho^*(1 + \epsilon)$ and $\rho/(1 + \epsilon)$)
- Probability $1 - \delta$ that estimate is within error bound ϵ (absolute/relative)
- For absolute error, relation between δ , ϵ and M is given by

$$M \geq \frac{\ln(2/\delta)}{2\epsilon^2}.$$

- Example: set $\delta = .05$, $\epsilon = .1$, $M \geq \ln(40/2) / .01 = 185$; change ϵ to $.01$, $M = 18,500$ etc

- For relative error:
$$M \geq 3 \frac{\ln(2/\delta)}{P(\mathbf{y})\epsilon^2}.$$

- Note: number of samples depends inversely on the actual value of $P(\mathbf{y})$

- Note: we don't actually know $P(\mathbf{y})$ so hard to estimate M

Rejection Sampling

- To compute conditional probabilities $P(\mathbf{Y}=\mathbf{y} \mid \mathbf{E}=\mathbf{e})$, we compute $\#[\mathbf{Y}=\mathbf{y} \mid \mathbf{E}=\mathbf{e}] / \#[\mathbf{E}=\mathbf{e}]$
- We throw away samples where $\mathbf{E} \neq \mathbf{e}$
- If $P(\mathbf{e})$ is small, we will throw away a large fraction of the samples; getting enough useful samples for reliable estimation will require a very large number of total samples
- Intuitive Solution: Set variables in \mathbf{E} to have value \mathbf{e} , then generate samples
 - Problem: generated samples are biased, see next slide

Likelihood Weighting

- We can just set evidence variables to the observed values to avoid rejecting samples
- However, this biases the samples. Suppose that evidence is $S=s^1$;
- Apply normal forward sampling, we would start from prior distributions of D and I , so $I=i^1$ in only 30% of the cases; this is clearly not consistent with $S=s^1$.
- In rejection sampling, if we started with $I=i^1$ then $S=s^1$ in 80% of the samples (from CPD); if we started with $I=i^0$ then $S=s^1$ in only 5% samples
- Weight samples by their likelihood: (i^1, s^1) by .8, (i^0, s^1) by .05
 - This weighting compensation is intuitive, formal justification comes later
 - Generalizes to case where more than one variable is evidence variable (product of probabilities for generating evidence variables given their parents; specifics in Algorithm 12.2, next slide).

$$\hat{P}_{\mathcal{D}}(\mathbf{y} \mid e) = \frac{\sum_{m=1}^M w[m] \mathbf{I}\{\mathbf{y}[m] = \mathbf{y}\}}{\sum_{m=1}^M w[m]}$$

LW Algorithm

Algorithm 12.2 Likelihood-weighted particle generation

Procedure LW-Sample (
 \mathcal{B} , // Bayesian network over \mathcal{X}
 $Z = z$ // Event in the network
)

- 1 Let X_1, \dots, X_n be a topological ordering of \mathcal{X}
- 2 $w \leftarrow 1$
- 3 **for** $i = 1, \dots, n$
- 4 $u_i \leftarrow x \langle \text{Pa}_{X_i} \rangle$ // Assignment to Pa_{X_i} in x_1, \dots, x_{i-1}
- 5 **if** $X_i \notin Z$ **then**
- 6 Sample x_i from $P(X_i \mid u_i)$
- 7 **else**
- 8 $x_i \leftarrow z \langle X_i \rangle$ // Assignment to X_i in z
- 9 $w \leftarrow w \cdot P(x_i \mid u_i)$ // Multiply weight by probability of desired value
- 10 **return** $(x_1, \dots, x_n), w$

Next Class

- Read sections 12.2 and 12.3 of the KF book